

15-740/18-740 Computer Architecture, Fall 2011

Midterm Exam I KEY

Instructor: Onur Mutlu

Teaching Assistants: Justin Meza, Yoongu Kim

Date: October 24, 2011

Name: _____

Problem I (50 points) :	_____
Problem II (12 points) :	_____
Problem III (30 points) :	_____
Problem IV (35 points) :	_____
Problem V (20 points) :	_____
Problem VI (15 points) :	_____
Bonus Problem VII (18 points) :	_____
Legibility and Name (5 points) :	_____
Total (167 + 18 points) :	_____

Instructions:

1. This is a closed book exam. You are allowed to have one letter-sized cheat sheet.
2. No electronic devices may be used.
3. This exam lasts 1 hour 50 minutes.
4. Clearly indicate your final answer.
5. Please show your work when needed.
6. Please write your initials on every odd page.
7. Please make sure that your answers to all questions (and all supporting work that is required) are contained in the space provided.

- **Be cognizant of time.** Do not spend too much time on one question.
- **Be concise.** You will be penalized for verbosity.
- **Show work when needed.** You will receive partial credit.
- **Write legibly.** Show your final answer.

I. Potpourri (50 points)

1. Static vs. Dynamic Scheduling

In class, we discussed that the compiler can reorder instructions dynamically to minimize the effect of data dependencies on execution time. However, it is limited in its ability to do so. Give three reasons as to why this is the case (each reason < 15 words):

Unknown branch direction.

Unknown memory addresses.

Unknown cache hit/miss behavior.

2. MIPS

The original MIPS architecture design did not have hardware interlocks between pipeline stages. Hence the name MIPS, Microprocessor without Interlocking Pipeline Stages. Assume a four-stage pipelined implementation of MIPS, with the pipeline stages we discussed in class: F, D, E, SR. Assume no data forwarding between stages.

```
add R1 <- R1, R2
add R1 <- R1, R3
add R1 <- R1, R4
add R1 <- R1, R5
```

a) Whose task is it to guarantee that the above piece of code executes correctly? (< 10 words.)

The compiler.

b) How so? Explain. Give the resulting code if there need to be changes to the code.

The compiler can insert NOPs in between dependent instructions:

```
add R1 <- R1, R2
NOP
add R1 <- R1, R3
NOP
add R1 <- R1, R4
NOP
add R1 <- R1, R5
```

3. Addressing Modes

a) What is the benefit of the autoincrement addressing mode? (< 15 words.)

Reduces code size by combining a memory instruction and address calculation instruction.

b) What kind of data structures and access patterns would benefit from this mode? (< 30 words.)

Array traversal.

4. CRAY

CRAY-1 was a vector processor. Yet, it had the fastest scalar unit of all processors at the time it was designed. Explain why the designers made this choice. (< 20 words.)

Amdahl's Law: Speedup of Cray was limited by sequential bottlenecks.

5. Fine-Grained Multithreading

In lecture, we discussed the idea of fine-grained multithreading as a technique to tolerate latency.

a) Which of the following latencies can fine-grained multithreading tolerate in a pipelined processor?
Explain very briefly (< 20 words).

Control-flow dependencies (Please circle YES or NO.)

YES NO

Branch is resolved by the time the next instruction is fetched from the same thread.

Data dependencies (Please circle YES or NO.)

YES NO

Dependencies will be resolved by the time the next instruction is fetched from the same thread.

Resource contention latency (Please circle YES or NO.)

YES NO

In the strictest definition of fine-grained multithreading, all latencies can be tolerated with enough threads.

b) Does fine-grained multithreading have any benefit in terms of hardware complexity?

YES NO

Explain why so (< 20 words).

No hardware interlocks, no data forwarding, no branch prediction logic.

c) Does fine-grained multithreading have any downside in terms of hardware complexity?

YES NO

Explain why so (< 20 words).

Multiple register files (hardware contexts) needed.

6. Interconnect

a) Company X is designing a multi-core processor with 100 cores. Their initial design connects the cores with a single shared bus. Give two reasons as to why this is not a good idea. (< 20 words each.)

Limited bandwidth → high contention.

Cannot operate at high frequency due to electrical loading.

b) An engineer suggests using a crossbar to connect each core to every other, instead. Give two reasons as to why this is a good idea. (< 20 words each.)

High bandwidth → low contention.

Can operate at high frequency.

Give two reasons as to why this is not a good idea. (< 20 words each.)

A lot more hardware than a single bus: $O(n^2)$ versus $O(1)$.

Power; hard to scale to a larger number of cores.

7. Exceptions

Consider the following state of the reorder buffer in an out-of-order execution engine:

ROB-ID	OP	COMPLETE?	EXCEPTION?
1	ld	NO	NO
2	add	NO	NO
3	br	NO	NO
4	add	YES	NO
5	ld	YES	YES
6	mul	NO	NO
7	ld	NO	NO

Assume the load in ROB-ID 5 generated an access protection exception because it references a NULL pointer. Give two reasons why this could be the case (< 20 words each):

Program has a bug.

Load in ROB-ID 5 is on wrong path.

8. Hybrid Main Memory

Justin gave a lecture on some recent work in hybrid main memory data placement policies. A hybrid main memory consists of a small amount of DRAM which acts as a cache to a large amount of non-volatile memory (NVM), and Justin discussed the trade offs of storing certain types of data on either device. Lets say you run a program on a hybrid main memory system and observe that (1) the total number of memory accesses is large and (2) the DRAM cache miss rate is large (meaning not many accesses are serviced from DRAM), but the program is hardly slowed down. What characteristics could the program have for this to occur (be specific)? (< 20 words.)

High row-buffer locality (in PCM).

9. Accelerated Critical Sections

Justin and Yoongu run their favorite parallel application, ResearchIdeaGenerator, on the same machine with 16 symmetric cores with one difference: it either executes each critical sections on the core that encounters the critical section or it executes every critical section on one core that is dedicated to the execution of critical sections and serial bottlenecks. They measure the performance of ResearchIdeaGenerator and find that the first machine completes the application in 240 minutes whereas the second in 300 minutes. They dig deeper into the problem and find that the cache miss rate is significantly higher in the second machine than in the first and that is the cause of the performance degradation. Provide two reasons as to why this could be the case (< 20 words each):

Shared data working set does not fit in a single cache but it does in multiple caches.

The amount of private data to be shipped to the single core can be larger than the amount of shared data that stays in one cache.

II. Branch Misprediction Penalty (12 points)

Assume a baseline in-order fetch, in-order dispatch processor design has four pipeline stages corresponding to fetch, decode, execute, and writeback; only executes arithmetic and branch instructions; and all instructions only require one execute pipeline stage. Assume you are benchmarking the performance of a program which is composed of 80% arithmetic instructions and 20% branch instructions.

- a) What is the misprediction penalty (in number of cycles) for this processor design?

3 cycles

- b) Calculate the throughput of the baseline processor design in terms of retired (written back) instructions per cycle (IPC) assuming no branches are mispredicted by the branch predictor.

1 IPC

- c) Calculate the IPC assuming 50% of the branches are mispredicted by the branch predictor.

Consider 10 retired instructions: 8 were arithmetic and 2 were branches. Of those 2 branches, 1 was predicted correctly and 1 was predicted incorrectly. The 1 which was predicted incorrectly led to 3 cycles of additional work. So, $IPC = \boxed{10/13}$. (If 2 assumed in part a, 10/12; if 4 assumed in part a, 10/14.)

One alternative design might be to change the fetch stage into two stages as well as change the decode stage into two stages (for a total of six stages in the new pipeline) in order to increase frequency.

- d) What misprediction rate would your branch predictor on this new six-stage pipeline design need to have in order to achieve the same IPC as the baseline four-stage pipeline design (still assuming the baseline design utilizes a 50% accurate branch predictor)?

Notice that the misprediction penalty times the misprediction rate should be equal between the two designs. So, we must solve for x in the equation $3(0.5) = 5x \Rightarrow x = \boxed{0.3}$. (If 2 assumed in part a, 0.25; if 4 assumed in part a, 0.33.)

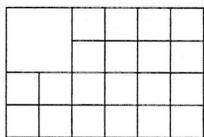
III. Amdahl's Law (30 points)

Let us assume that you are designing a multi-core processor to be fabricated on a fixed silicon die area budget of A . As the architect, you are to partition this total area of A into one “strong” core and many “weak” cores. The strong core will have an area of S , while the weak cores will each have an area of 1 (but there will be $A - S$ of them). Assume that the single-thread performance of a core scales with the square root of its area.

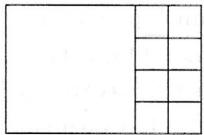
1. “Strong” vs. “Weak”

On this multiprocessor, we will execute a workload where P fraction of its work is *infinitely* parallelizable, and where $1 - P$ of its work is serial.

- Configuration X: $A = 24, S = 4$ (One strong core, 20 weak cores.)



- Configuration Y: $A = 24, S = 16$ (One strong core, eight weak cores.)



Assume that the serial portion of the workload executes only on the strong core and the parallel portion of the workload executes on both the strong and weak cores.

- a) What is the speedup of the workload on configuration X? (Compared to the execution time on a single-core processor of area 1.)

$$\left(\frac{1-P}{2} + \frac{P}{2+20} \right)^{-1}$$

- b) What is the speedup of the workload on configuration Y? (Compared to the execution time on a single-core processor of area 1.)

$$\left(\frac{1-P}{4} + \frac{P}{4+8} \right)^{-1}$$

c) For what range of P does the workload run faster on Y than X?

$$\left(\frac{1-P}{4} + \frac{P}{4+8}\right)^{-1} > \left(\frac{1-P}{2} + \frac{P}{2+20}\right)^{-1}$$

$$\frac{1-P}{4} + \frac{P}{12} > \frac{1-P}{2} + \frac{P}{22}$$

$$P\left(\frac{1}{12} - \frac{1}{22}\right) < \frac{1-P}{4}$$

$$P\left(\frac{1}{6} - \frac{1}{11}\right) < \frac{1-P}{2}$$

$$P\frac{11-6}{66} < \frac{1-P}{2}$$

$$P\frac{5}{66} < \frac{1}{2} - \frac{P}{2}$$

$$P\frac{38}{66} < \frac{1}{2}$$

$$P < \frac{33}{38}$$

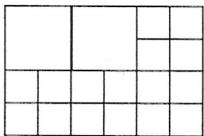
d) For workloads that have limited parallelism, which configuration would you recommend and why? (< 20 words.)

Y because of its better speedup.

2. Dual-Core Execution

In class, we discussed some ways in which multiple cores can be utilized to emulate the single-thread performance of a larger core. For this question, we assume that two strong cores can operate in a “dual-core execution” manner to achieve the single-thread performance of an even “stronger” core that is $1.5\times$ faster than any single strong core.

- Configuration DCE: $A = 24, S_1 = 4, S_2 = 4$ (Two strong cores, sixteen weak cores.)



Assume that dual-core execution occurs only during the serial portion of the workload and that during the parallel portion, the two strong cores separate from each other (on-the-fly) and operate as two independent cores. Furthermore, the serial portion executes only on the “dual-core”, and the parallel portion executes on all the cores.

- a) What is the speedup of the workload on configuration DCE? (Compared to the execution time on a single-core processor of area 1.)

$$\left(\frac{1-P}{3} + \frac{P}{2+2+16}\right)^{-1}$$

b) For what range of P does the workload run faster on DCE than X?

$$\left(\frac{1-P}{3} + \frac{P}{20}\right)^{-1} < \left(\frac{1-P}{2} + \frac{P}{22}\right)^{-1}$$

$$\frac{1-P}{3} + \frac{P}{20} > \frac{1-P}{2} + \frac{P}{22}$$

$$P \left(\frac{1}{20} - \frac{1}{22} \right) < \frac{1-P}{6}$$

$$P \left(\frac{1}{10} - \frac{1}{11} \right) < \frac{1-P}{3}$$

$$P \left(\frac{11-10}{110} \right) < \frac{1}{3} - \frac{P}{3}$$

$$P \left(\frac{1}{110} \right) + \frac{P}{3} < \frac{1}{3}$$

$$P \left(\frac{3}{330} + \frac{110}{330} \right) < \frac{1}{3}$$

$$P \left(\frac{113}{110} \right) < 1$$

$$P < \frac{110}{113}$$

c) Suppose you are executing workloads with a wide range of parallelism. Which configuration would you choose? DCE or Y? Why?

DCE as it offers good performance for a wide range of P .

IV. Reorder Buffer & Register Alias Table (35 points)

In class and the review session, we covered the operation of an out-of-order execution engine that employs a reorder buffer (ROB), a consolidated physical register file, frontend and backend (retirement) register alias tables, and checkpointing. The figure below shows the state of these structures after the last multiply instruction at ROB-ID 10 is fetched and renamed. Only the relevant fields of each entry of each structure are shown.

Reorder Buffer							Backend RAT	
ROB-ID	Valid	OP	Dest. ArchReg	Dest. PhyReg	Completed	PhyReg		
1	1	ld	r1	p55	0	r1	p1	
2	1	add	r1	p57	0	r2	p2	
3	1	mul	r3	p61	1	r3	p3	
4	1	br	-	-	0	r4	p4	
5	1	add	r1	p80	1	r5	p5	
6	1	mul	r5	p90	0	r6	p6	
7	1	br	-	-	0	r7	p7	
8	1	add	r3	p100	0	r8	p8	
9	1	br	-	-	0			
10	1	mul	r7	p88	0			

RAT1		RAT2		RAT3		RAT4	
Rdy	PhyReg	Rdy	PhyReg	Rdy	PhyReg	Rdy	PhyReg
r1	p80	1	p80	0	p57	1	p80
r2	p2	1	p2	1	p2	1	p2
r3	p100	0	p100	1	p61	1	p61
r4	p4	1	p4	1	p4	1	p4
r5	p90	0	p90	1	p5	0	p90
r6	p6	1	p6	1	p6	1	p6
r7	p7	0	p88	1	p7	1	p7
r8	p8	1	p8	1	p8	1	p8

CP3	Frontend RAT	CP1	CP2
-----	--------------	-----	-----

The figure was drawn by a computer architect who was tracing the operation of a program. Unfortunately, his playful kid erased parts of the state using whiteout (marked with gray boxes in your handout). Your task: complete the entire state by filling in the gray boxes. First, start with answering the following questions.

- a) How many architectural registers does the depicted machine have?

8

Four register alias tables (RATs) are depicted in the frontend. These four correspond to the current “Frontend RAT” which is used for renaming incoming instructions as well as three checkpointed RATs. A checkpoint of the frontend RAT was taken when each of the branches was renamed (hence, three checkpoints since there are three branches in the ROB). Remember that each entry in each RAT contains a “Ready (R)” bit (0 indicates “not ready”; as you recall, this bit is also called the “Valid” bit) and a tag corresponding to the physical register ID.

b) What is the purpose of checkpointing the frontend RAT after a branch is renamed? (< 20 words.)

Fast state recovery on misprediction.

There is also a backend (retirement) RAT depicted, but it does not contain any "Ready" bits.

c) Why does the backend RAT not contain a ready bit for each register? (< 20 words.)

Because it represents architectural state which is always valid.

We have not told you which of the RATs correspond to the "Frontend RAT" and the three checkpoints. Given all you know about the state of the machine, your task is to identify which RAT corresponds to these.

d) In the box under each RAT, please identify the RAT as one of the following:

- "Frontend RAT" (current)
- "CP1" (Checkpointed RAT corresponding to the first branch - Branch at ROB-ID 4)
- "CP2" (Checkpointed RAT corresponding to the second branch - Branch at ROB-ID 7)
- "CP3" (Checkpointed RAT corresponding to the third branch - Branch at ROB-ID 9)

Note that the RATs' order of appearance on paper does not correspond to any timeline.

e) Now, go ahead and fill all the gray boxes in the figure. (If you make a mess, clearly indicate your final answer. Use blank sheets if needed.)

f) Show the state of the Backend RAT after the third branch (at ROB-ID 9) is retired (assuming all branches were predicted correctly and there are no exceptions).

Backend RAT

PhyReg
r1
r2
r3
r4
r5
r6
r7
r8

(The state of the Backend RAT is the same as CP3, without the valid bits.)

Now, assume instead that the second branch (at ROB-ID 7) was mispredicted. Also assume its misprediction is resolved long before the load is completed and the multiply at ROB-ID 6 is dependent on the load. Immediately after the branch is resolved, the state is recovered and the processor starts fetching from the correct path.

- g) Show the state of the Frontend RAT immediately after misprediction recovery and before the first instruction on the correct path is renamed.

Frontend RAT	
	Rdy PhyReg
r1	1 p80
r2	1 p2
r3	1 p61
r4	1 p4
r5	0 p90
r6	1 p6
r7	1 p7
r8	1 p8

(The state of the Frontend RAT is the same as CP2, and R5 is not ready because the multiply writing to it is not yet complete.)

- h) Given the above information, can you construct the dataflow graph of the code that is currently in the ROB? If so, draw the dataflow graph below. If not, why not? Tell us what other structures' state would you need to see in the machine to be able to construct the data flow graph. (< 20 words.)

No, because we do not know the source registers of instructions. Need reservation stations.

V. Tomasulo's Algorithm (20 points)

For this question, we'll consider an in-order fetch, out-of-order dispatch processor with fetch, decode, execute, reorder, and writeback stages, and a future register file. Our system has two functional units: an adder that has three pipeline stages, and a multiplier that has five pipeline stages. Assume that tags and data are broadcast on the common data bus during the last cycle of the execute stage. During cycle 0, we fetch the first of six instructions on our processor. We show the state of the future file at cycle 0 below as well as the state of the machine at cycle 9. ("—" is used to indicate information that is unknown.)

Future register file at cycle 0

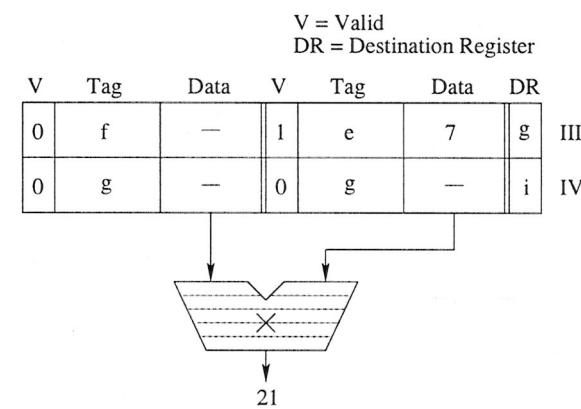
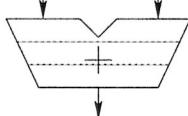
	V	Tag	Data
r0	1	c	3
r1	1	b	5
r2	1	z	16
r3	1	x	0
r4	1	y	100
r5	1	d	-1
r6	1	a	2

Future register file at cycle 9

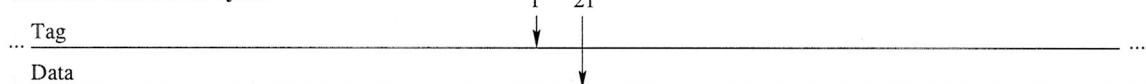
	V	Tag	Data
r0	1	c	3
r1	1	b	5
r2	0	f	-
r3	1	e	7
r4	0	g	-
r5	1	d	-1
r6	0	j	-

Reservation stations at cycle 9

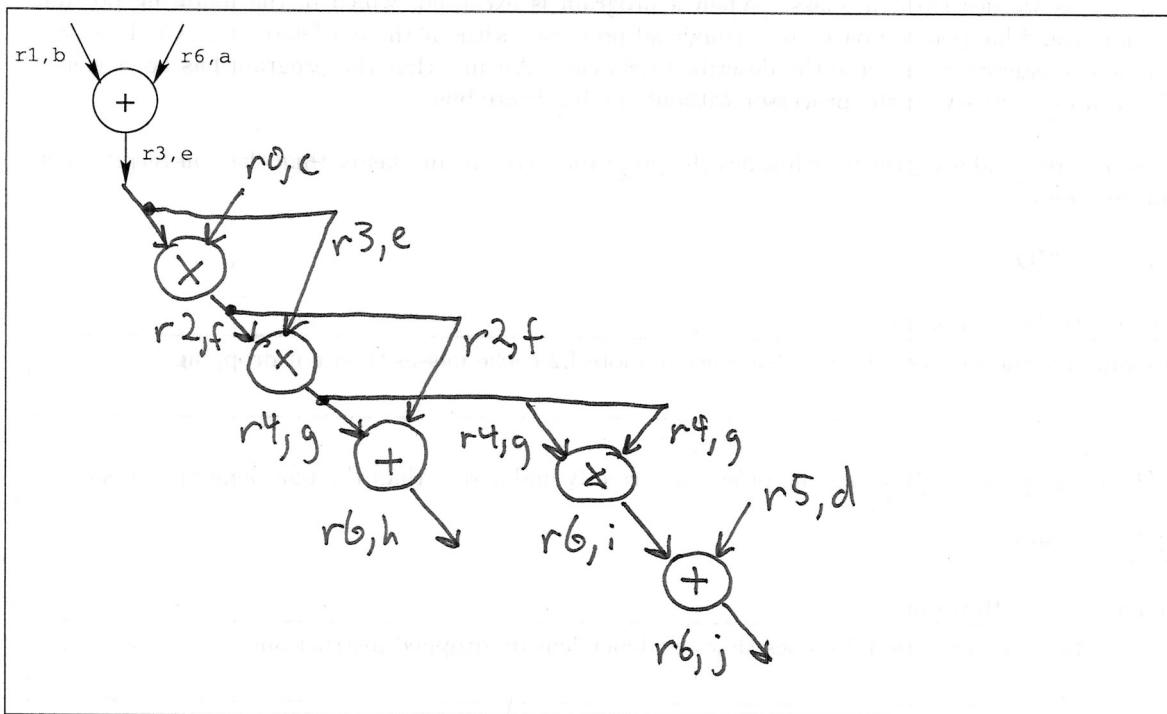
	V	Tag	Data	V	Tag	Data	DR
I	0	f	-	0	g	-	h
II	0	i	-	1	d	-1	j



Common data bus at cycle 9



- a) Your first task is to use only the supplied information to draw the data flow graph for the six instructions which have been fetched. Label nodes with the operation being performed and edges with the architectural register and corresponding tag being used. Note that the first instruction (which was fetched in cycle 0) was ADD $r_3 \leftarrow r_1, r_6$, and this instruction has already committed its data to architectural state.



- b) Then, use the data flow graph to fill in the table below with the six instructions being executed on the processor. When referring to registers, please use their architectural names (r_0 through r_6) with source registers written in ascending order (as shown in the filled-in entry below). Note that the operation for instruction 5 has already been filled in.

OP = operation
 Dest. = destination architectural register
 Src.1 = source architectural register 1
 Src.2 = source architectural register 2

	OP	Dest.	Src.1	Src.2
Instruction 1	ADD	r_3	r_1	r_6
Instruction 2	MUL	r_2	r_0	r_3
Instruction 3	MUL	r_4	r_2	r_3
Instruction 4	ADD	r_6	r_2	r_4
Instruction 5	MUL	r_6	r_4	r_4
Instruction 6	ADD	r_6	r_5	r_6

VI. Runahead Potpourri I (15 points)

1. Runahead and Hardware Bugs

A runahead processor was designed with an unintended hardware bug: every other instruction in runahead mode is dropped by the processor after fetch. All other behavior of the runahead mode is exactly as we described in class. When a program is executed, which of the following possible behaviors could happen compared to a runahead processor without the hardware bug? Circle YES if there is a possibility to observe the described behavior. Assume that the program has no bug in it and executes correctly on the processor without the hardware bug.

- a) The buggy runahead processor finishes the program correctly and faster than the non-buggy runahead processor.

YES NO

Explain why (< 16 words).

Dropping instructions enables the discovery of more L2 cache misses than not dropping.

- b) The buggy processor finishes the program correctly and slower than the non-buggy processor.

YES NO

Explain why (< 16 words).

Not able to generate some L2 misses that are dependent on dropped instructions.

- c) The buggy processor executes the program incorrectly.

YES NO

Explain why (< 16 words).

Not possible because runahead mode cannot change program behavior as it is purely speculative.

2. Runahead vs. Multithreading

Others have suggested an alternative to runahead execution is to use multithreading of the following form: have multiple hardware contexts in hardware (just like in fine-grained multithreading) and when one incurs an L2 cache miss, start executing from the other context. This method is commonly called switch-on-L2-miss multithreading, or more generally switch-on-event (SoE) multithreading.

- a) What are two advantages of runahead execution over SoE multithreading? (< 20 words each.)

Improves single thread performance.

Does not require switching in of a separate thread context.

- b) What are two advantages of SoE over runahead execution? (< 20 words each.)

Performs work that is not going to be discarded for sure.

Doesn't require propagation of invalid bits.

VII. BONUS: Runahead Potpourri II (18 points)

1. Runahead vs. a Larger Instruction Window

In our performance simulation experiments, we found that a processor with a 128-entry instruction window and runahead execution achieves significantly higher performance than a processor with a 1024-entry window but no runahead execution, for a particular program A.

a) Explain why this could be the case. Provide as many reasons as possible, but no more than the correct ones. Enumerate each reason as 1., 2., 3., ...

1. Runahead runs further than 1024 instructions and discovers cache misses a 1024-entry window cannot uncover.
2. Runahead can run further than 1024 instructions and can train the branch predictor and prefetchers.
3. Larger instruction window processor can have a deeper pipeline or a slower clock frequency due to the larger structures associated with it.

For program B, we observed the opposite to be true: the 1024-entry window machine without runahead achieves significantly higher performance than a 128-entry window runahead machine.

b) Explain why this could be the case. Provide as many reasons as possible, but no more than the correct ones. Enumerate each reason as 1., 2., 3., ...

1. Runahead cannot find any cache misses (useless) and flushes the window for no reason.
2. 1024-entry window can tolerate latencies of non-L2 miss instructions better.
3. Runahead has flush penalty, 1024-entry window does not.
4. Runahead runs further than what 1024 entry window would allow but causes pollution and wrong prefetches.
5. Runahead mistrains the branch predictor and prefetchers.
6. There are many dependent misses in the program.

2. Runahead and Value Prediction

Recall that in conventional runahead execution, when a load instruction misses in the L2 cache, its destination register is marked as invalid (INV). INV bits are propagated through instruction execution in such a way that all registers that are dependent on L2 misses are marked INV.

a) What is the rationale for marking a register as INV? (< 20 words.)

So that non-existent values are not used for prefetching and branch resolution.

A clever engineer thinks of an alternative mechanism. Her suggestion is to assume that the value in the destination register of an L2 miss load is zero, instead of marking the register INV: in other words, predict the data value of an L2 miss load to be zero.

b) What are two advantages of this proposal over conventional runahead execution? (< 20 words each.)

No INV bits needed → simpler.

Can potentially resolve mispredicted INV branches correctly now.

c) When and why can this proposal be beneficial for performance (or some other metric) compared to conventional runahead? Provide a code example if it helps clarity.

Can keep runahead on a useful prefetch path. E.g., can get out of a loop that keeps getting mispredicted due to unavailable data, or linked list traversal.

d) What are two disadvantages of the engineer's proposal over conventional runahead execution? (< 20 words each.)

Can lead to more branch mispredictions in runahead mode.

Can lead to higher energy consumption and slower progress during runahead mode because otherwise-INV instructions are now all executed, even though this may not lead to any benefit.

e) When and why can this proposal degrade performance (or some other metric) compared to conventional runahead? Provide a code example if it helps clarity.

Can push runahead on wrong path by turning a correctly-predicted branch to an incorrectly-resolved one.

SCRATCH PAD

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

When you have finished working on this page, fold it in half and place it in your portfolio.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

When you have finished working on this page, fold it in half and place it in your portfolio.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.

Use this page for scratch work. You may use this page for calculations, drawing, or any other work you need to do.