


What is a bus?

- **Slow vehicle that many people ride together**
 - well, true...
- **A bunch of wires...**

MicroComputer Engineering Bus slide 1

A Bus is:


- a shared communication link
- a single set of wires used to connect multiple subsystems



- **a Bus is also a fundamental tool for composing large, complex systems**
 - systematic means of abstraction

MicroComputer Engineering Bus slide 2


Advantages of Buses



- **Versatility:**
 - New devices can be added easily
 - Peripherals can be moved between computer systems that use the same bus standard
- **Low Cost:**
 - A single set of wires is shared in multiple ways
- **Manage complexity by partitioning the design**

MicroComputer Engineering Bus slide 3


Disadvantage of Buses



- **It creates a communication bottleneck**
 - The bandwidth of that bus can limit the maximum I/O throughput
- **The maximum bus speed is largely limited by:**
 - The **length** of the bus
 - The **number** of devices on the bus
 - The need to support a range of devices with:
 - Widely varying latencies
 - Widely varying data transfer rates

MicroComputer Engineering Bus slide 4

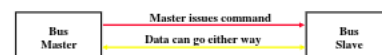
The General Organization of a Bus



- **Control lines:**
 - Signal requests and acknowledgments
 - Indicate what type of information is on the data lines
- **Data lines** carry information between the source and the destination:
 - Data and Addresses
 - Complex commands

MicroComputer Engineering Bus slide 5

Master versus Slave



- **A bus transaction** includes two parts:
 - Issuing the command (and address) – request
 - Transferring the data – action
- **Master is the one who starts the bus transaction by:**
 - issuing the command (and address)
- **Slave is the one who responds to the address by:**
 - Sending data to the master if the master ask for data
 - Receiving data from the master if the master wants to send data

MicroComputer Engineering Bus slide 6

- **Processor-Memory Bus (design specific)**

- Short and high speed
- Only need to match the memory system
 - Maximize memory-to-processor bandwidth
- Connects directly to the processor
- Optimized for cache block transfers
- **Backplane Bus (standard or proprietary)**
 - Backplane: an interconnection structure within the chassis
 - Allow processors, memory, and I/O devices to coexist
 - Cost advantage: one bus for all components
- **I/O Bus (industry standard)**
 - Usually is lengthy and slower
 - Need to match a wide range of I/O devices
 - Connects to the processor-memory bus or backplane bus

Processor/Memory Bus




Diagram illustrating a shared bus architecture. A central horizontal line represents the **Backplane Bus**. On the left, a box labeled **Processor** is connected to the bus. On the right, a box labeled **Memory** is connected to the bus. Below the bus, four cylinder icons represent **I/O Devices**, each connected to the bus by a vertical line.

- **A single bus (the backplane bus) is used for:**
 - Processor to memory communication
 - Communication between I/O devices and memory
- **Advantages: Simple and low cost**
- **Disadvantages: slow and the bus can become a major bottleneck**
- **Example: IBM PC - AT**

The diagram illustrates a shared bus architecture. A central horizontal line represents the 'Processor Memory Bus'. On the left, a box labeled 'Processor' is connected to the bus. On the right, a box labeled 'Memory' is connected to the bus. Three 'Bus Adaptor' ovals are connected to the bus. Each adaptor is connected to an 'I/O Bus' (vertical line) which then connects to two disk-like storage units.

- I/O buses tap into the processor-memory bus via bus adaptors:
 - Processor-memory bus: mainly for processor-memory traffic
 - I/O buses: provide expansion slots for I/O devices
- Apple Macintosh-II
 - NuBus: Processor, memory, and a few selected I/O devices
 - SCSI Bus: the rest of the I/O devices


The diagram illustrates a multi-processor system architecture. At the top, a red line represents the 'Processor Memory Bus' connecting a 'Processor' block on the left to a 'Memory' block on the right. Below the processor, a blue line represents the 'Backplane Bus'. This bus connects to two 'Bus Adaptor' blocks (represented by ovals). Each 'Bus Adaptor' is connected to a yellow 'I/O Bus' (represented by a horizontal line). Each 'I/O Bus' is then connected to two disk-like storage units (represented by cylinders). The entire system is enclosed in a red border.

- **A small number of backplane buses tap into the processor-memory bus**
 - Processor-memory bus is used for processor memory traffic
 - I/O buses are connected to the backplane bus
- **Advantage: loading on the processor bus is greatly reduced**

Processor-Memory Bus


■ This bus connects the CPU to RAM

- Designed for maximal bandwidth
- Usually wide, 32 bits or more
- To further increase bandwidth we use a Cache
- Burst access between cache and memory, **early restart**




MicroComputer Engineering Bus slide 13

Memory Bus




MicroComputer Engineering Bus slide 14

On Chip Cache (1st level)



MicroComputer Engineering Bus slide 15

What defines a bus?



MicroComputer Engineering Bus slide 16

Synchronous and Asynchronous Bus

■ Synchronous bus:

- Includes a clock in the control lines
- A fixed protocol for communication that is relative to the clock
- Advantage: involves very little logic and can run very fast
- Disadvantages:
 - Every device on the bus must run at the same clock rate
 - To avoid clock skew, they cannot be long if they are fast

■ Asynchronous bus:

- It is not clocked
- It can accommodate a wide range of devices
- It can be lengthened without worrying about clock skew
- It requires a handshaking protocol

MicroComputer Engineering Bus slide 17

Synchronous Bus

■ A single clock controls the protocol

- Pros
 - Simple (one FSM)
 - Fast
- Cons
 - Clock skew limits bus length
 - All devices work on the same speed (clock)

■ Suitable for Processor-Memory Bus

MicroComputer Engineering Bus slide 18

Asynchronous Bus


■ Uses a **handshaking** to implement a transaction protocol

- Pros
 - Versatile, generic protocols
 - Dynamic data rate
- Cons
 - More complex, two communication FSMs
 - Slower, (but usually can be made quite fast)

MicroComputer Engineering Bus slide 19

Asynchronous Protocol

- ReadReq
- DataReady
- Ack



MicroComputer Engineering Bus slide 20

Busses so far



Bus Master: has ability to control the bus, initiates transaction

Bus Slave: module activated by the transaction

Bus Communication Protocol: specification of sequence of events and timing requirements in transferring information.

Asynchronous Bus Transfers: control lines (req, ack) serve to orchestrate sequencing.

Synchronous Bus Transfers: sequence relative to common clock.

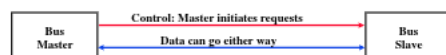
MicroComputer Engineering Bus slide 21

Bus Transaction

- Arbitration
- Request
- Action

MicroComputer Engineering Bus slide 22

Arbitration: Obtaining Access to the Bus



- One of the most important issues in bus design:
 - How is the bus reserved by a devices that wishes to use it?
- Chaos is avoided by a master-slave arrangement:
 - Only the bus master can control access to the bus:
 - It initiates and controls all bus requests
 - A slave responds to read and write requests
- The simplest system:
 - Processor is the only bus master
 - All bus requests must be controlled by the processor
 - Major drawback: the processor is involved in every transaction

MicroComputer Engineering Bus slide 23

Bus Arbitration

- Bus Master, (initiator usually the CPU)
- Slave, (usually the Memory)

Arbitration signals

- BusRequest
- BusGrant
- BusPriority
 - Higher priority served first
 - Fairness, no request is locked out

MicroComputer Engineering Bus slide 24

Multiple Potential Bus Masters: the Need for Arbitration

- **Bus arbitration scheme:**
 - A bus master wanting to use the bus asserts the bus request
 - A bus master cannot use the bus until its request is granted
 - A bus master must signal to the arbiter after finish using the bus
- **Bus arbitration schemes usually try to balance two factors:**
 - Bus priority: the highest priority device should be serviced first
 - Fairness: Even the lowest priority device should never be completely locked out from the bus


MicroComputer Engineering Bus slide 25

Bus Arbitration Schemes

- **Daisy chain arbitration**
 - Grant line runs through all device, highest priority device first.
 - Single device with all request lines.
- **Centralized**
 - Many request/grant lines
 - Complex controller, may be a bottleneck
- **Self Selection**
 - Many request lines
 - The one with highest priority self decides to take bus
- **Collision Detection (Ethernet)**
 - One request line
 - Try to access bus,
 - If collision device backoff
 - Try again in random + exponential time

MicroComputer Engineering Bus slide 26


The Daisy Chain Bus Arbitrations Scheme



- **Advantage: simple**
- **Disadvantages:**
 - Cannot assure **fairness**: A low-priority device may be locked out indefinitely
 - The use of the daisy chain grant signal also limits the bus speed

MicroComputer Engineering Bus slide 27

Centralized Parallel Arbitration



- **Used in essentially all processor-memory busses and in high-speed I/O busses**

MicroComputer Engineering Bus slide 28


Simplest bus paradigm



- All agents operate synchronously
- All can source / sink data at same rate
- => simple protocol
 - just manage the source and target

MicroComputer Engineering Bus slide 29


Simple Synchronous Protocol



- **Even memory busses are more complex than this**
 - memory (slave) may take time to respond
 - it need to control data rate

MicroComputer Engineering Bus slide 30

Typical Synchronous Protocol



- Slave indicates when it is prepared for data xfer
- Actual transfer goes at bus rate

MicroComputer Engineering Bus slide 31

Increasing the Bus Bandwidth


- **Separate versus multiplexed address and data lines:**
 - Address and data can be transmitted in one bus cycle if separate address and data lines are available
 - Cost: (a) more bus lines, (b) increased complexity
- **Data bus width:**
 - By increasing the width of the data bus, transfers of multiple words require fewer bus cycles
 - Example: SPARCstation 20's memory bus is 128 bit wide
 - Cost: more bus lines
- **Block transfers:**
 - Allow the bus to transfer multiple words in back-to-back bus cycles
 - Only one address needs to be sent at the beginning
 - The bus is not released until the last word is transferred
 - Cost: (a) increased complexity, (b) increased response time (latency) for request

MicroComputer Engineering Bus slide 32

Pipelined Bus Protocols

Attempt to initiate next address phase during current data phase

Single master example (proc-to-cache)



MicroComputer Engineering Bus slide 33

Increasing Transaction Rate on Multimaster Bus

- **Overlapped arbitration**
 - perform arbitration for next transaction during current transaction
- **Bus parking**
 - master can hold onto bus and performs multiple transactions as long as no other master makes request
- **Overlapped address / data phases (prev. slide)**
 - requires one of the above techniques
- **Split-phase transaction bus (Command queueing)**
 - completely separate address and data phases
 - arbitrate separately for each
 - address phase yield a tag which is matched with data phase
- "All of the above" in most modern mem busses

MicroComputer Engineering Bus slide 34

Split Transaction Protocol

When we don't need the bus, release it!

- Improves throughput
- Increases latency and complexity



MicroComputer Engineering Bus slide 35

The I/O Bus Problem

- **Designed to support wide variety of devices**
 - full set not known at design time
- **Allow data rate match between arbitrary speed devices**
 - fast processor - slow I/O
 - slow processor - fast I/O


MicroComputer Engineering Bus slide 36

High Speed I/O Bus

- Examples
 - Raid controllers
 - Graphics adapter
- Limited number of devices
- Data transfer bursts at full rate
- DMA transfers important
 - small controller spools stream of bytes to or from memory
- Either side may need to stall transfer
 - buffers fill up


MicroComputer Engineering Bus slide 37

Backplane Bus



MicroComputer Engineering Bus slide 38

Direct Memory Access (DMA)



MicroComputer Engineering Bus slide 39

Direct Memory Access

- **DMA Processor**
 - 1) Generates BusRequest, waits for Grant
 - 2) Put Address & Data on Bus
 - 3) Increase Address, back to 2 until finished
 - 4) Release Bus
- **Generates interrupt only**
 - When finished
 - If an error occurred

MicroComputer Engineering Bus slide 40

DMA and Virtual Memory

- If DMA uses Virtual address it needs to pass a TLB
 - Go through the CPU's TLB (no good)
 - A TLB in the DMA processor (needs updating)
- If DMA uses Physical address
 - Only transfer within one Page
 - We give the DMA a set of Physical addresses
 - (local TLB copy)

MicroComputer Engineering Bus slide 41

DMA and Caching

- **INCONSISTENCY problem**
 - We change the RAM contents, but not the cache
 - We write to HD but the RAM holds "old" information
- **Routing All DMA though CPU**
 - No good, spoils the idea!
- **Software handling of DMA**
 - Cache: Flush selected cache lines to RAM
 - Cache: Invalidate selected cache lines
 - TLB/OS: Do not allow access to these pages until DMA finished
- **Hardware Cache Coherence Protocol**
 - Complex, but very useful for multi processor systems

MicroComputer Engineering Bus slide 42

PCI Read/Write Transactions

- All signals sampled on rising edge
- Centralized Parallel Arbitration
 - overlapped with previous transaction
- All transfers are (unlimited) **bursts**
- Address phase starts by asserting **FRAME#**
- Next cycle "initiator" asserts cmd and address
- Data transfers happen on when
 - IRDY# asserted by master/initiator when ready to transfer data
 - TRDY# asserted by target when ready to transfer data
 - transfer when both asserted on rising edge
- **FRAME#** deasserted when master intends to complete only one more data transfer

MicroComputer Engineering Bus slide 43

PCI Read Transaction




Figure 3-1: Basic Read Operation

– Turn-around cycle on any signal driven by more than one agent

MicroComputer Engineering Bus slide 44

PCI Write Transaction




Figure 3-2: Basic Write Operation

MicroComputer Engineering Bus slide 45

PCI Optimizations

- Push bus efficiency toward 100% under common simple usage
 - like RISC
- Bus Parking
 - retain bus grant for previous master until another makes request
 - granted master can start next transfer without arbitration
- Arbitrary Burst length
 - initiator and target can exert flow control with xRDY
 - target can disconnect request with STOP (abort or retry)
 - master can disconnect by deasserting FRAME
 - arbiter can disconnect by deasserting GNT
- Delayed (pending, split-phase) transactions
 - free the bus after request to slow device

MicroComputer Engineering Bus slide 46

Additional PCI Issues

- Interrupts: support for controlling I/O devices
- Cache coherency:
 - support for I/O and multiprocessors
- Locks:
 - support timesharing, I/O, and MPs
- Configuration Address Space

MicroComputer Engineering Bus slide 47

Summary of Bus Options

- | | | |
|-----------------|--------------------------------------|------------------------------------|
| ■ Option | <i>High performance</i> | <i>Low cost</i> |
| ■ Bus width | Separate address & data lines | Multiplex address & data lines |
| ■ Data width | Wider is faster (e.g., 32 bits) | Narrower is cheaper (e.g., 8 bits) |
| ■ Transfer size | Multiple words has less bus overhead | Single-word transfer is simpler |
| ■ Bus masters | Multiple (requires arbitration) | Single master (no arbitration) |
| ■ Clocking | Synchronous | Asynchronous |
| ■ Protocol | pipelined | Serial |

MicroComputer Engineering Bus slide 48